

SST and Cycle-accurate Simulation of xBGAS

Jie Li
Ph.D. Student, TTU
05/03/2022

Overview


- ✧ SST Overview
- ✧ Simulating with SST
- ✧ SST for xBGAS
- ✧ Summary & Future Work

SST Overview


❖ Why SST?

- ❖ Many simulators are serial and unable to simulate very large systems.
 - ❖ Significant performance issues with tying many simulators together.
 - ❖ Modification on tightly-coupled simulations is difficult.
 - ❖ Difficult to simulate at different levels of accuracy.
- Lack of scalability
- Lack of flexibility


**The Structural Simulation Toolkit:
an open, modular, parallel, multi-objective,
multi-scale simulation framework for
scalability and flexibility.**



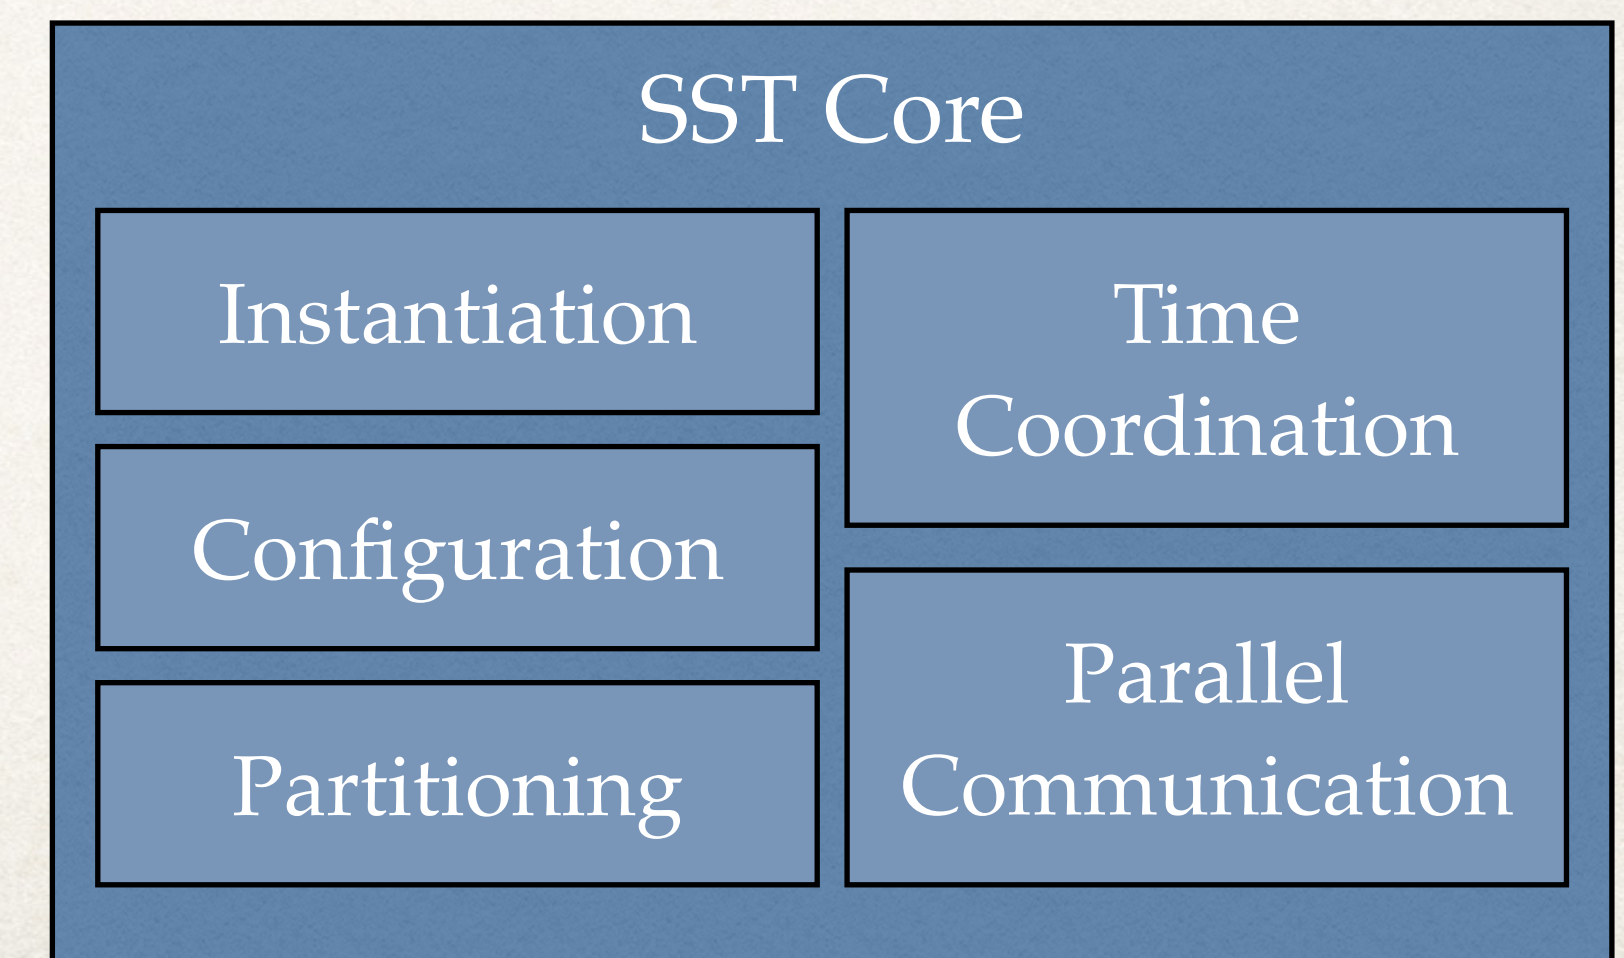
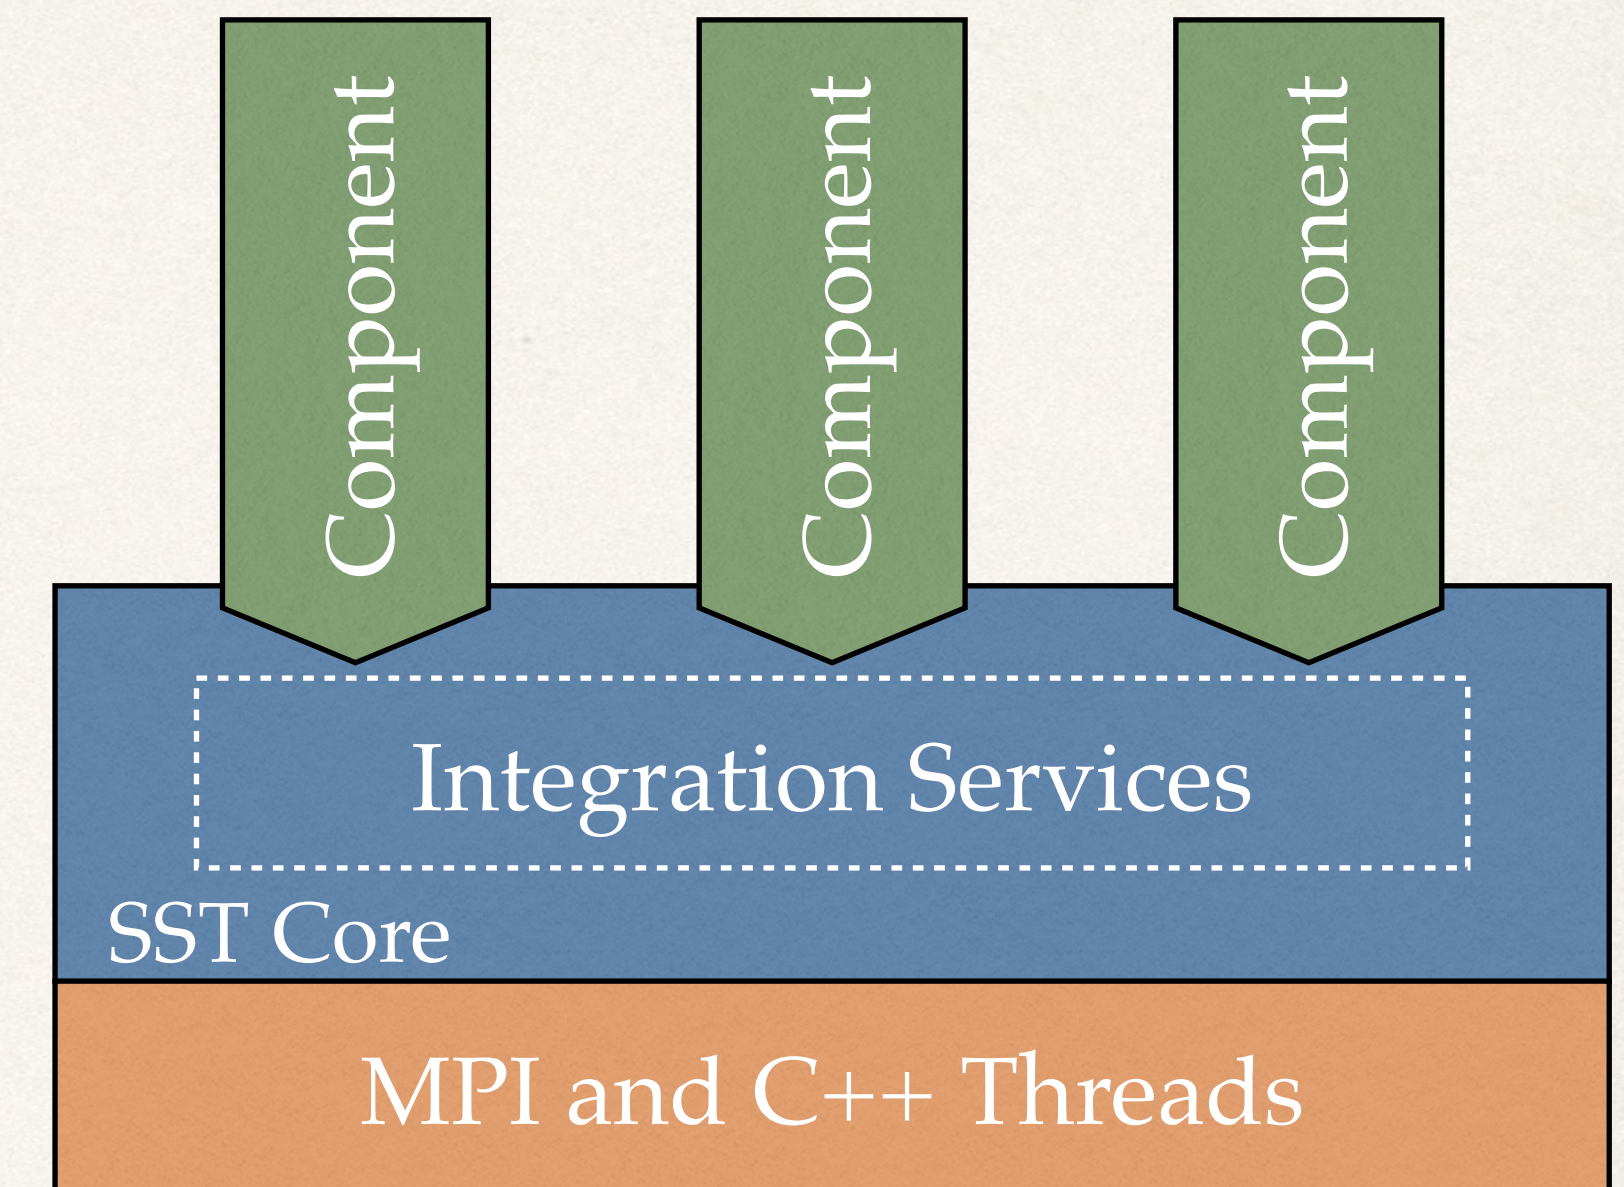
- ❖ Goals:
 - ❖ Create a standard architectural simulation framework for HPC
 - ❖ Enable “mix and match” of simulation components.
 - ❖ Custom architectures and custom tradeoff between accuracy and simulation time.
 - ❖ **Use supercomputers to design supercomputers.**
 - ❖ The simulation performance is optimized over MPI/Threads.



**Sandia
National
Laboratories**



- ❖ **SST Core** (framework):
 - ❖ The backbone of simulation.
 - ❖ Provides utilities and interfaces for simulation components.
 - ❖ Clocks, event exchange, statistics and parameter management, parallelism support, etc.
- ❖ **Components** (SST element libraries):
 - ❖ Components perform the actual simulation.
 - ❖ Built-in models include processors, memory, network, etc.
 - ❖ Compatible with many existing simulators: DRAMSim2, Spike, HMC-Sim, etc.



- ❖ **SST::Component**

- ❖ Simulation model

- ❖ **SST::Link**

- ❖ Communication path between two components.
- ❖ Poll or interrupt with EventHandler.

- ❖ **SST::Event**

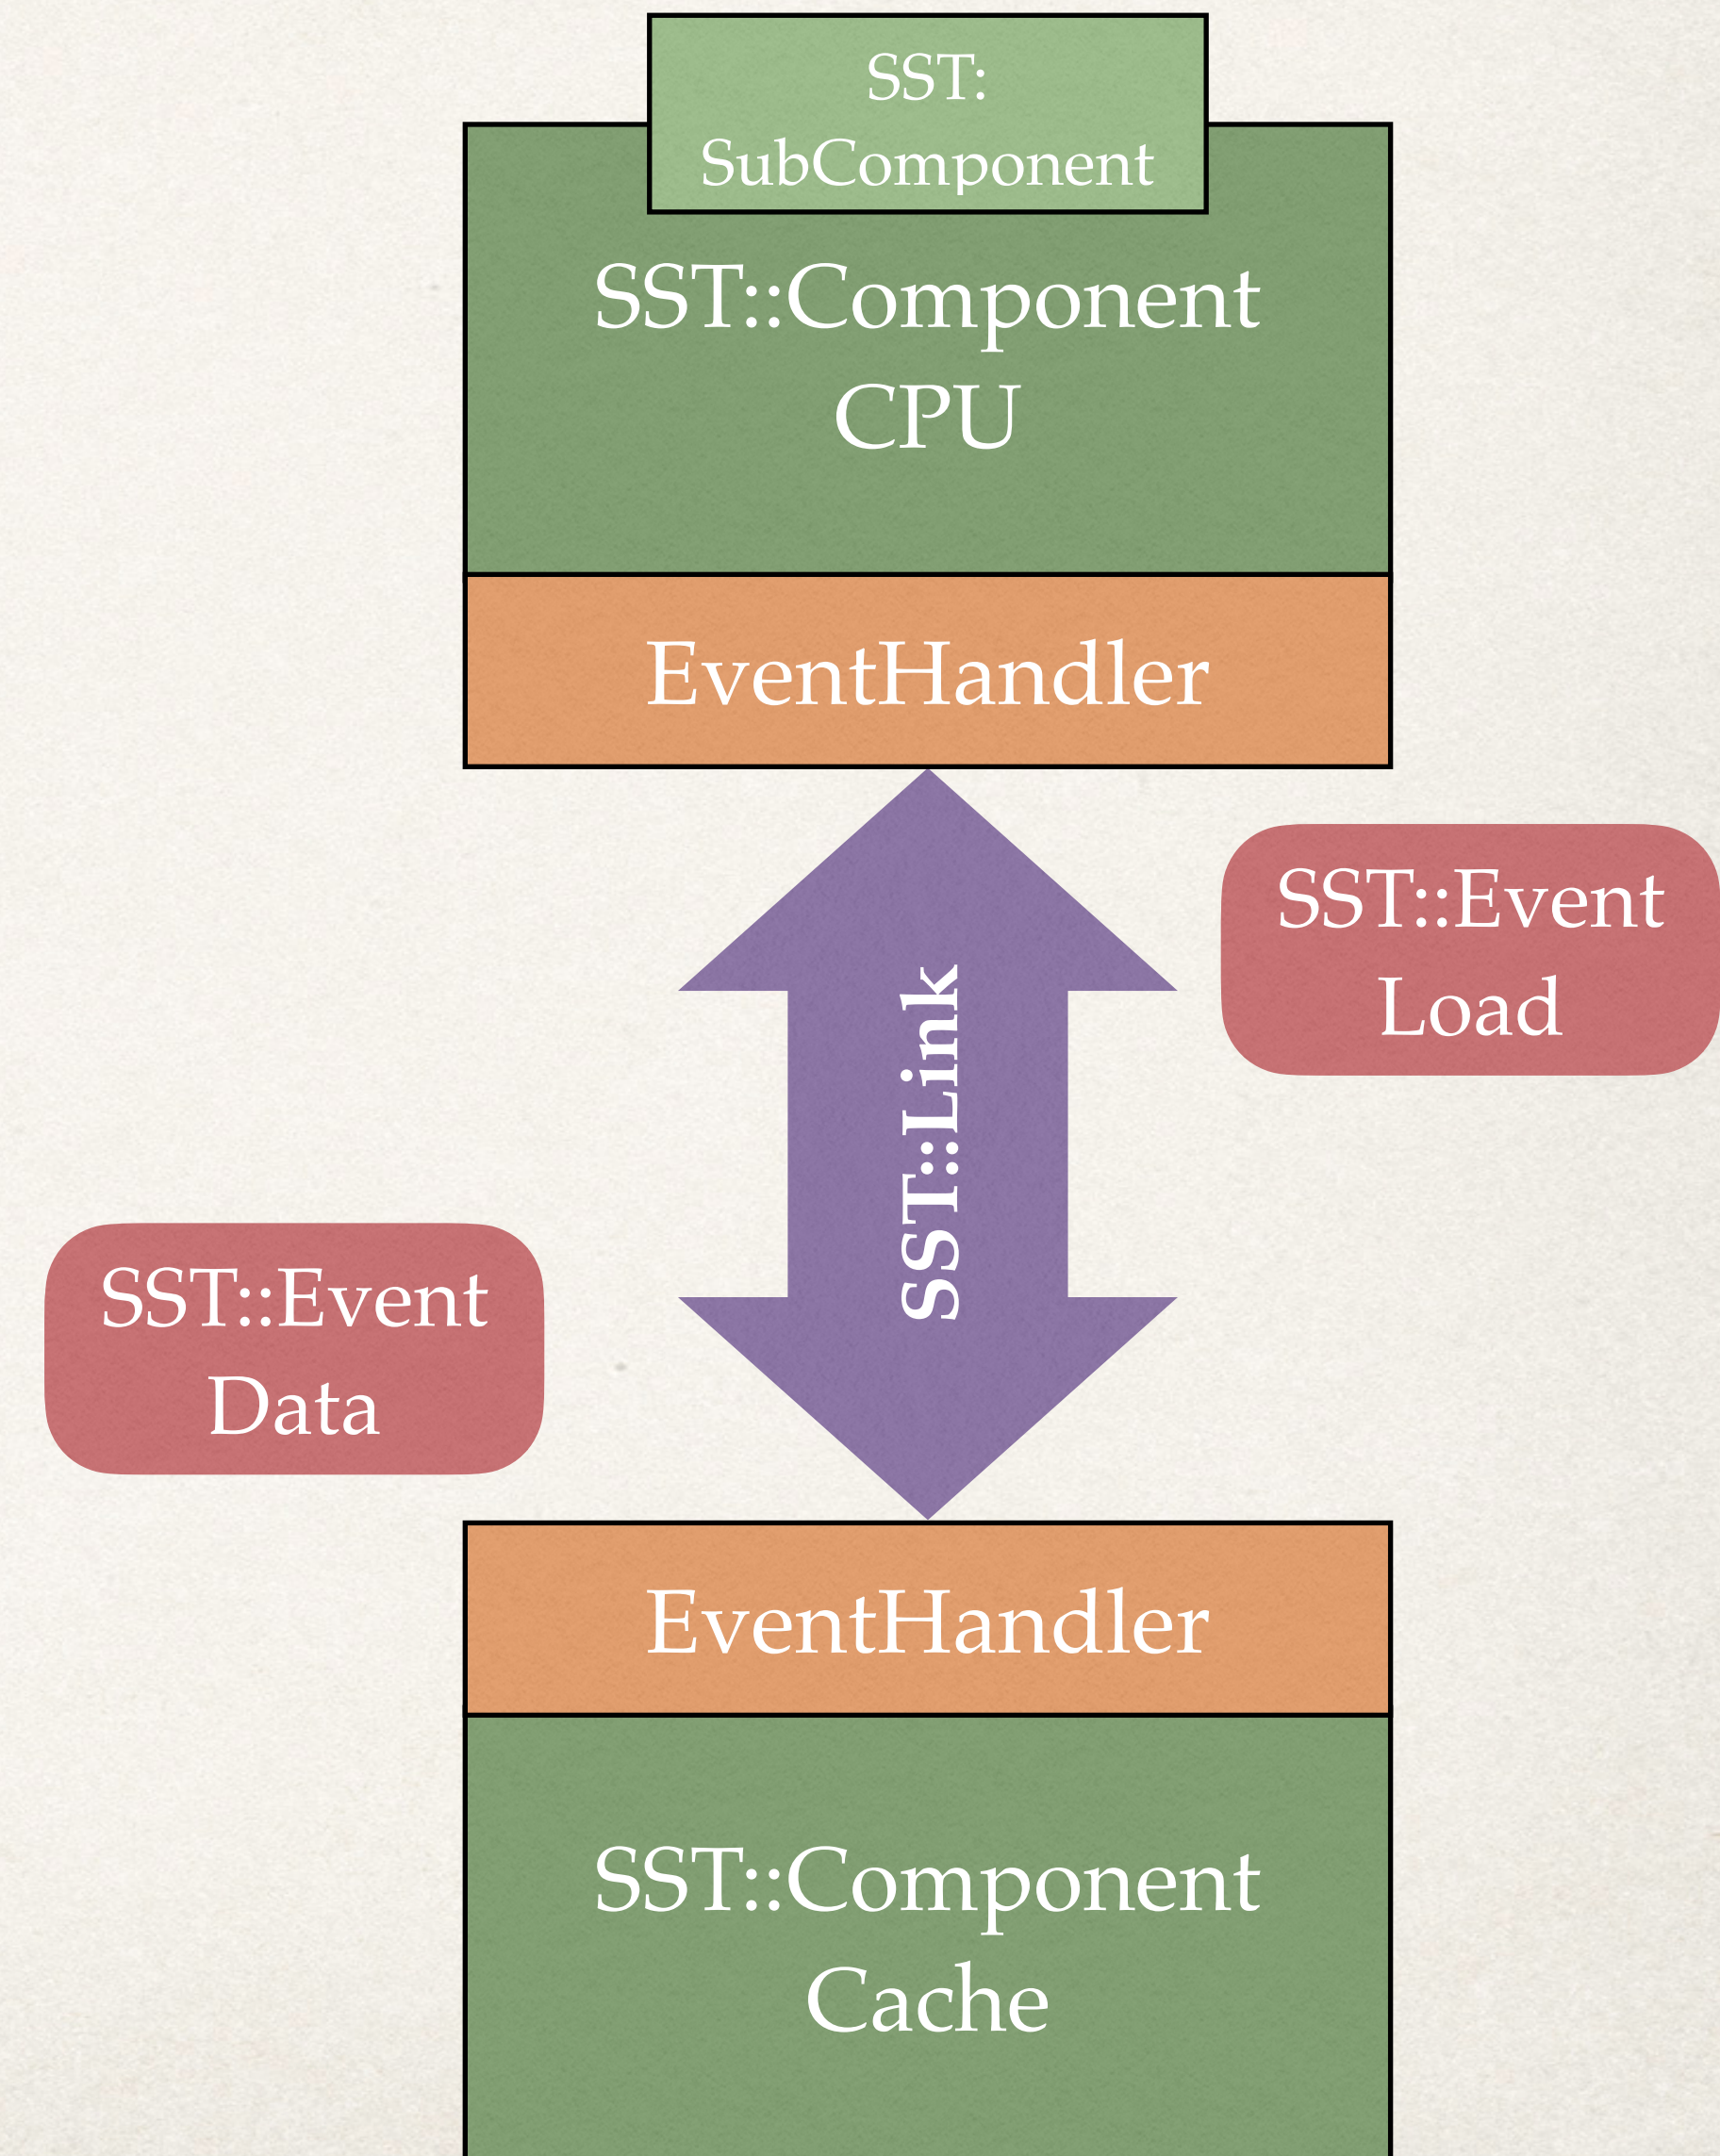
- ❖ A discrete event

- ❖ **SST::Clock::Handler**

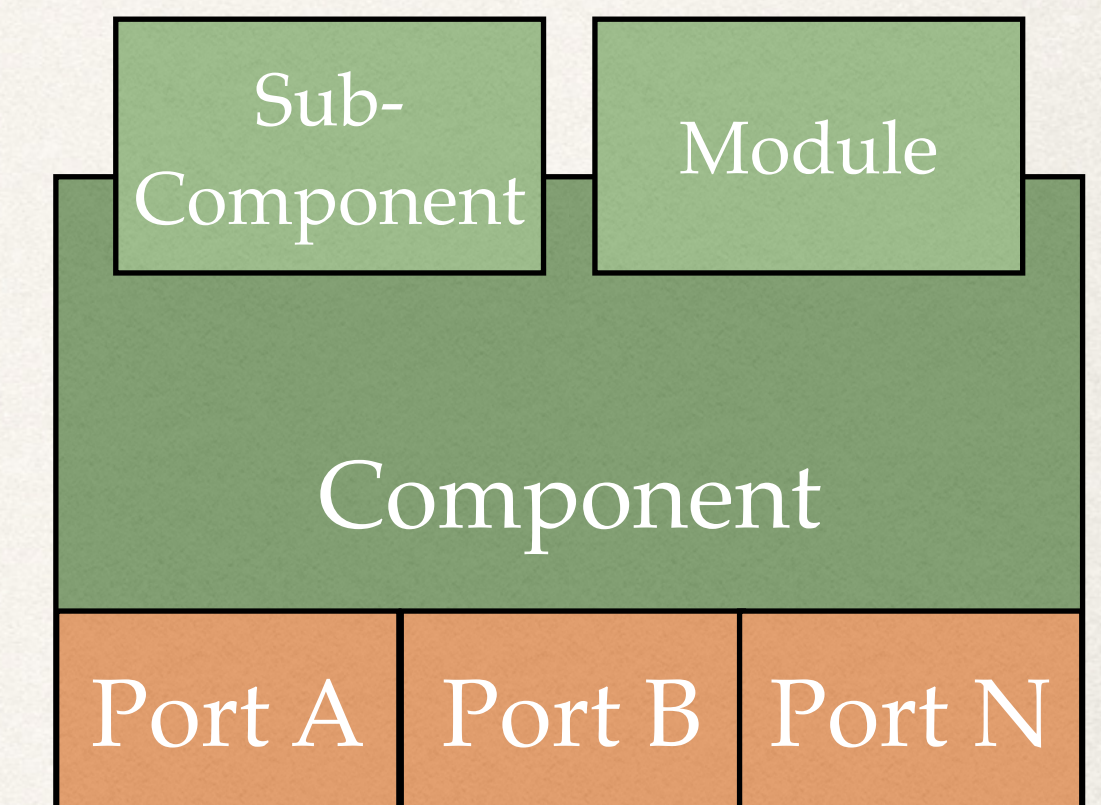
- ❖ Function to handle a clock tick.

- ❖ **SST::SubComponent**

- ❖ Add functionality to Components



- ❖ Basic building block of a simulation model. It performs the actual simulation.
- ❖ Use links and ports to communicate with other components.
 - ❖ Components define ports, links connect ports between components.
 - ❖ Polled: register a clock handler to poll the link.
 - ❖ Interrupt: register an event handler to be called when an event arrives.
 - ❖ Both: receive events on interrupt, send events on clock.
- ❖ Loads SubComponents and Modules for additional functionality.
 - ❖ Useful for parts of a component that should be swappable between simulations. E.g., scheduling algorithms, address mapping schemes, cache prefetchers.
 - ❖ Useful for exposing interfaces to other components. E.g., network provides a subcomponent that can be loaded to handle traffic into/out of network.



- ❖ Connects two components.
 - ❖ Connect a specific named “Port X” on component A to a “Port Y” on component B.
- ❖ The only mechanism by which components communicate.
 - ❖ Necessary for parallel simulation.
- ❖ Has minimum, non-zero latency for communication.
 - ❖ Except self-links.
 - ❖ Except during init/complete phases.
- ❖ Transparently handles any MPI/threaded communication.

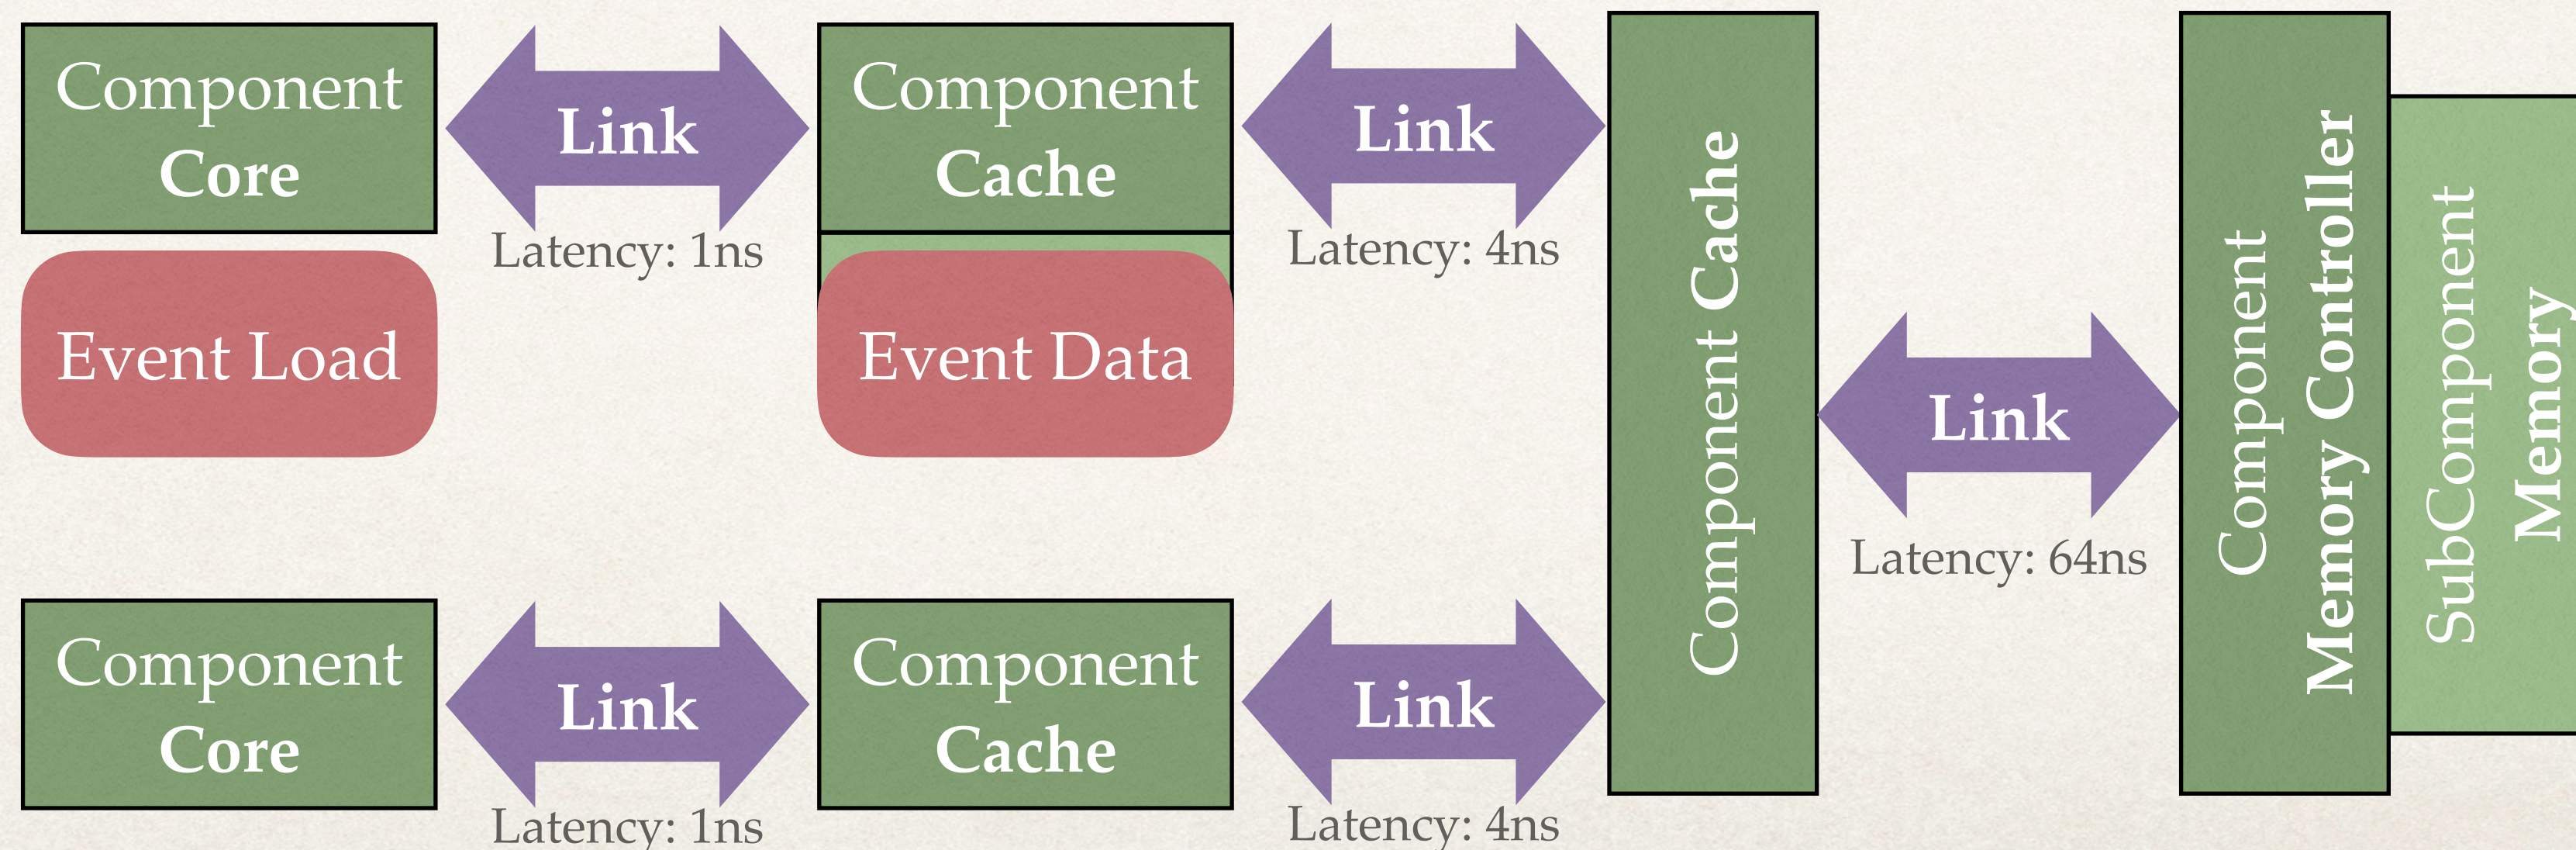


- ❖ Unit of communication between two components.
 - ❖ Packet format is up to the communicating components.
- ❖ Some standardized interfaces.
 - ❖ Facilitates “mix and match” capability.
 - ❖ sst/core/interfaces/
 - ❖ Memory (stdMem)
 - ❖ Defines commands & event format for communication with caches & memory.
 - ❖ Network (simpleNetwork)
 - ❖ Defines a header for events sent through a network component.

- ❖ Add additional functionality to a Component.
 - ❖ SubComponent (SC): shares common base class with Component (friend class of component).
 - ❖ Module (M): self-contained functionality (no access to Component class functionality).
- ❖ Provide modularity.
 - ❖ Generic interface which can be used by multiple SC/M.
 - ❖ Loaded by Components at runtime .
- ❖ Tightly-coupled with a component.
 - ❖ Components call SC/M as a C++ class instance.
 - ❖ Do not need to communicate via Links.
 - ❖ SC/M cannot exist by themselves.
- ❖ Example:
 - ❖ Miranda core loads a pattern generator subcomponent to generate a sequence of memory operations.

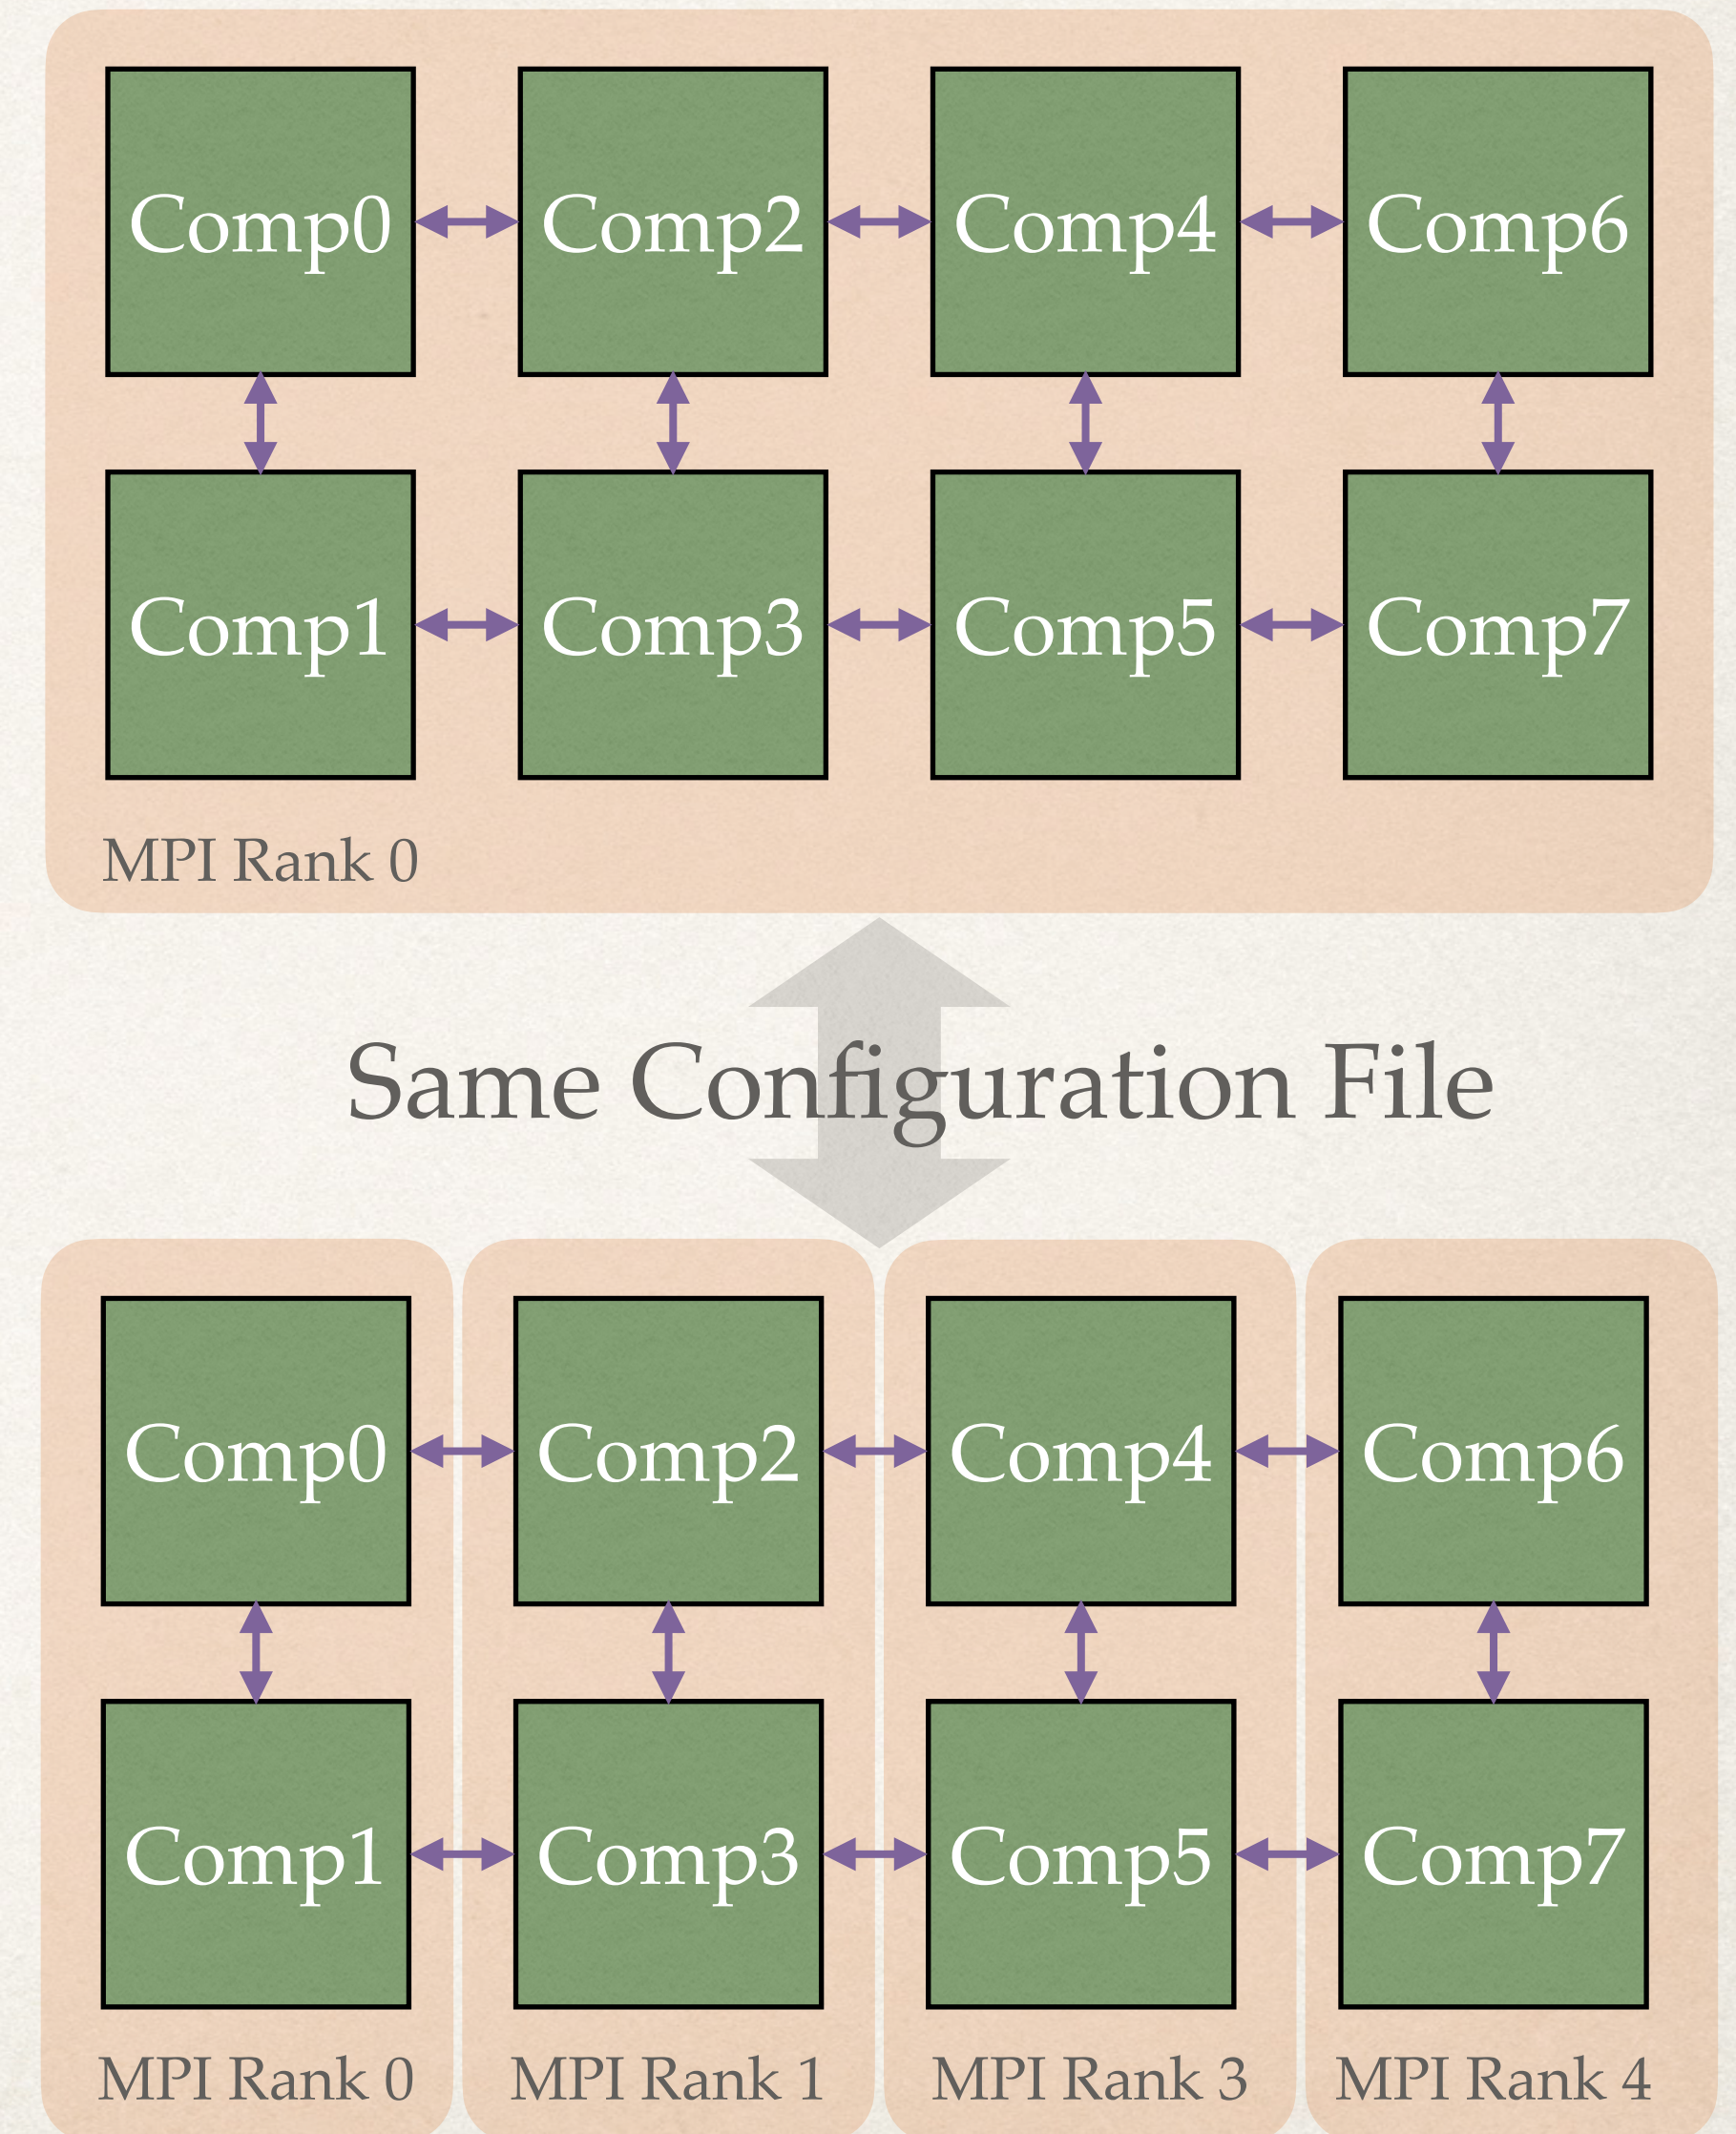


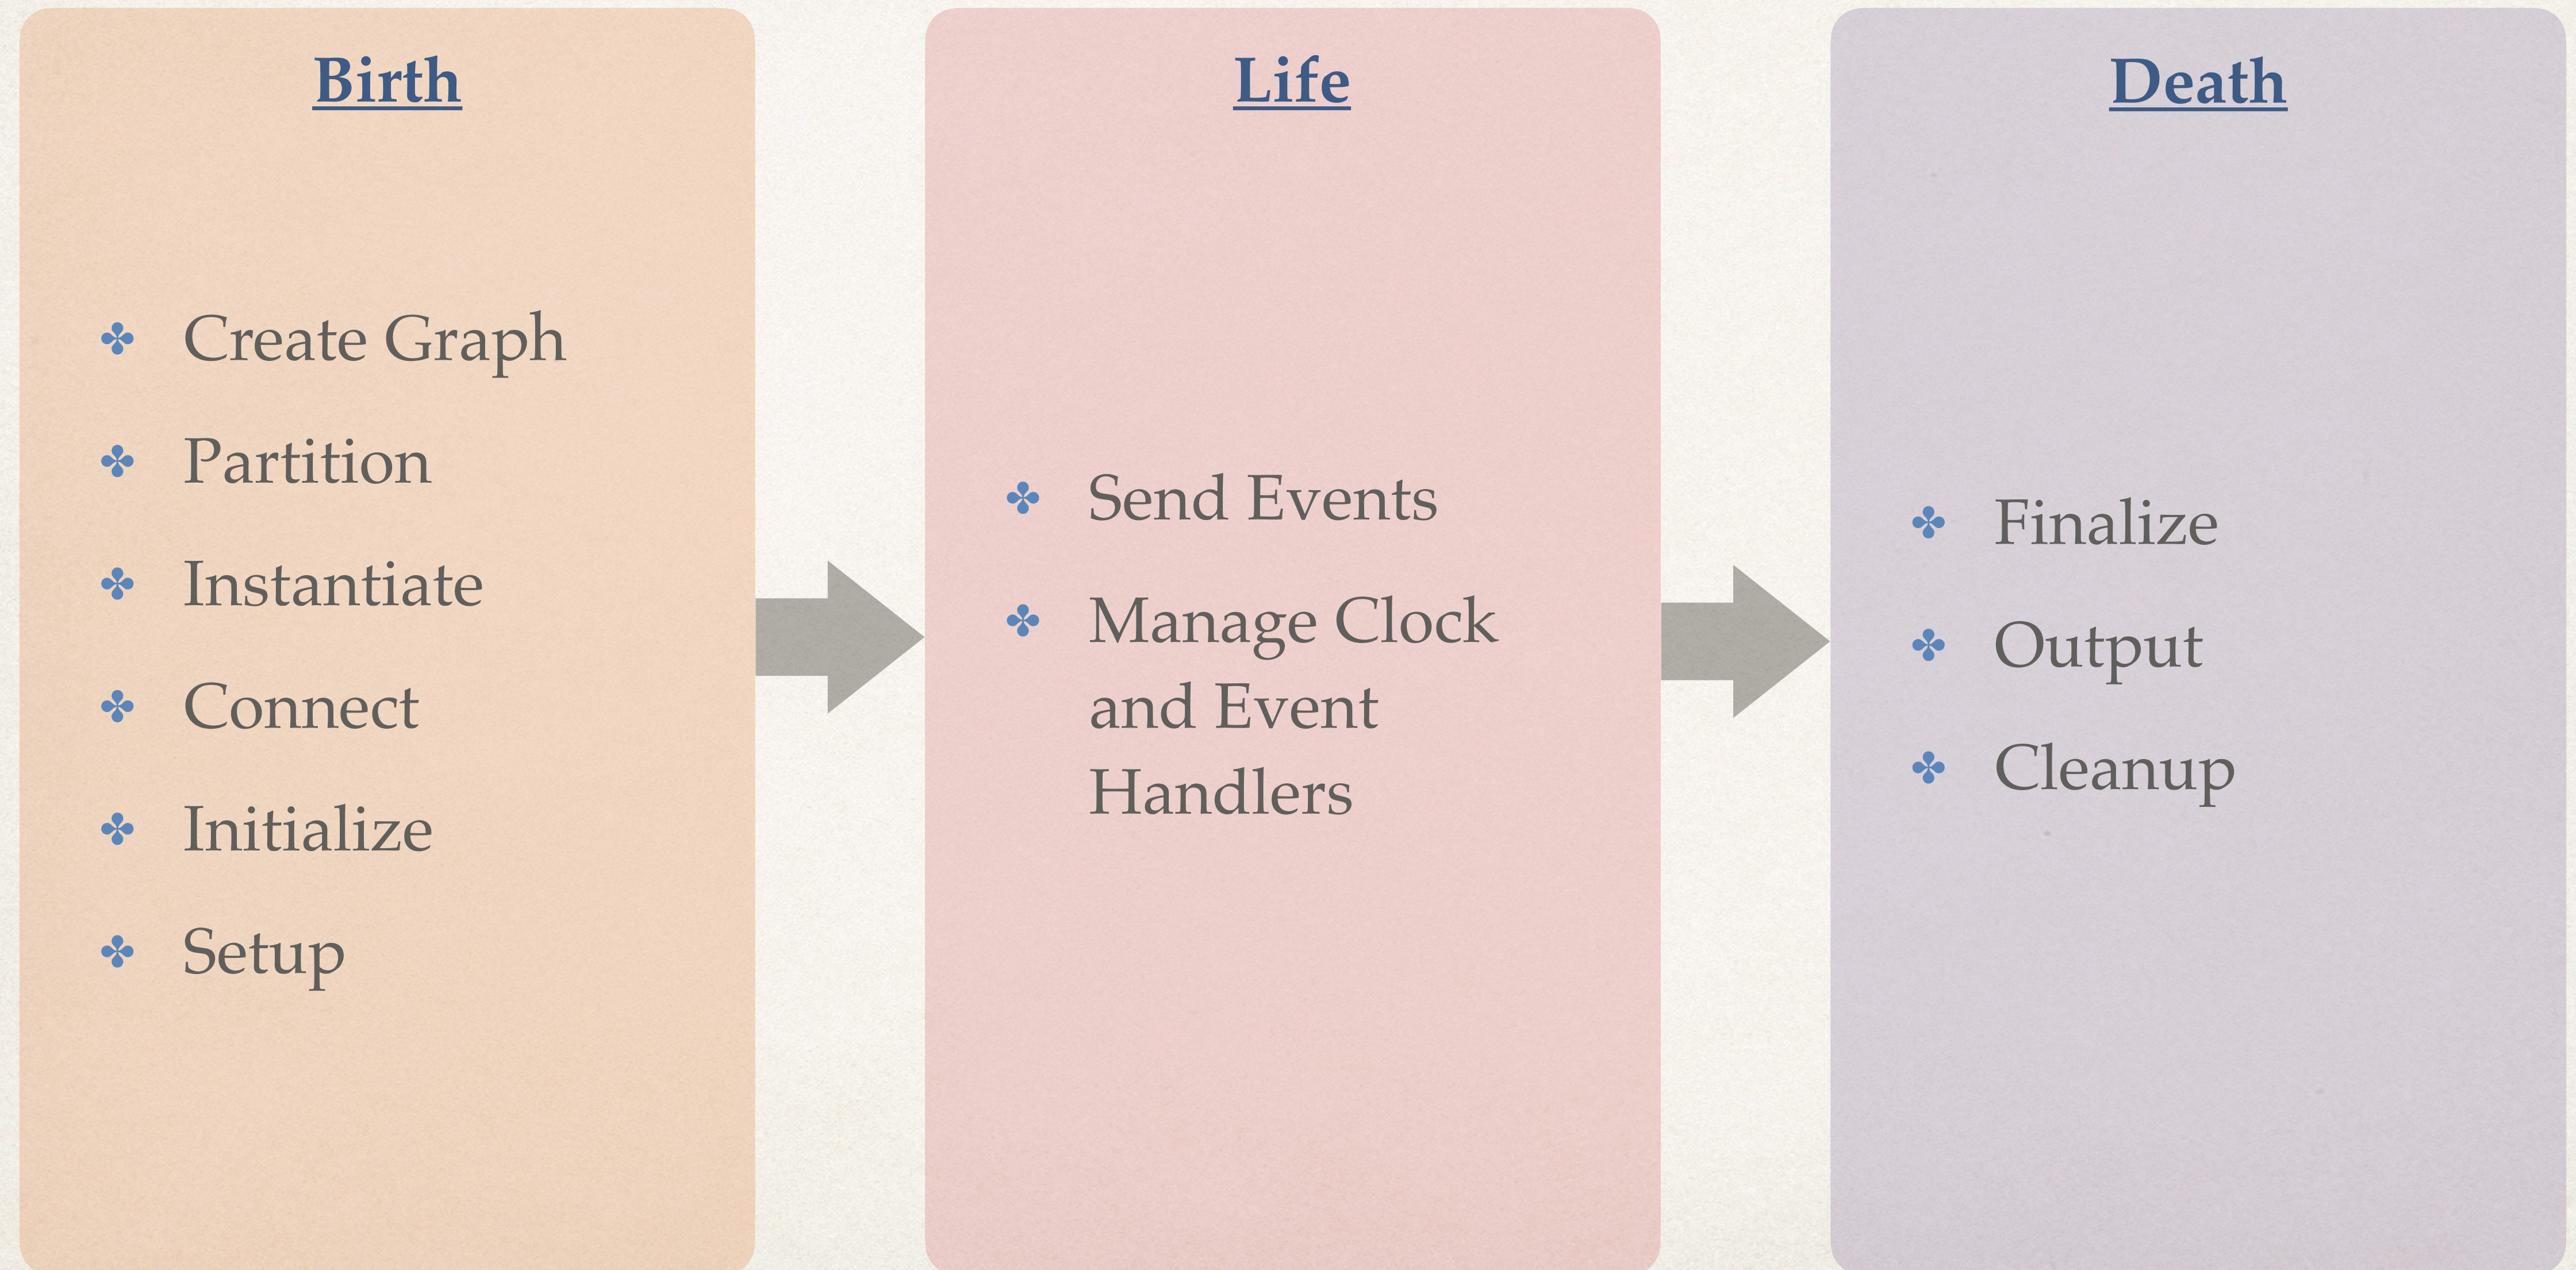
- ❖ SST simulations are comprised of **components** connected by **links**.
- ❖ Components communicate by sending **events** over links.
- ❖ Components define **ports** which are valid connection points for a link.
- ❖ Each link has a minimum latency.
- ❖ Components can use **subComponents** and **modules** for additional/ customizable functionality.



- ❖ SST was designed from the ground up to enable scalable, parallel simulations.
- ❖ Components are distributed among MPI ranks and / or threads.
- ❖ Links allow parallelism:
 - ❖ Components should communicate via links ONLY.
 - ❖ Transparently handle any MPI / thread communication.
 - ❖ Specified link-latency determines synchronization rate.


```
# Two ranks
$ mpirun -np 2 sst demo.py
# Two threads
$ sst -n 2 demo.py
# Two ranks with two threads
$ mirin -np 2 sst -n 2 demo.py
```





Simulating with SST

- ❖ Element libraries in the simulation:
 - ❖ **Miranda** - Simple core model that runs generated instruction streams.
 - ❖ **Generators** (SubComponents) produce memory access patterns, e.g., GUPSGenerator, RandomGenerator, STREAMBenchGenerator, etc.
 - ❖ **memHierarchy** - Various cache / memory system related subcomponents and modules.
 - ❖ **Cache** (Component) with coherence protocol SubComponent.
 - ❖ **Memory Controller** (Component) that loads a memory timing model (SubComponent).



SubComponent 7: STREAMBenchGenerator

Interface: SST::Miranda::RequestGenerator

NUM STATISTICS = 0

NUM PORTS = 0

NUM SUBCOMPONENT SLOTS = 0

NUM PARAMETERS = 7

PARAMETER 0 = verbose (Sets the verbosity output of the generator) [0]

PARAMETER 1 = n (Sets the number of elements in the STREAM arrays) [10000]

PARAMETER 2 = n_per_call (Sets the number of iterations to generate per call to the generation function) [1]

PARAMETER 3 = operandwidth (Sets the length of the request, default=8 (i.e. one double)) [8]

PARAMETER 4 = start_a (Sets the start address of the array a) [0]

PARAMETER 5 = start_b (Sets the start address of the array b) [1024]

PARAMETER 6 = start_c (Sets the start address of the array c) [2048]

STREAMBenchGenerator: Creates a representation of the STREAM benchmark

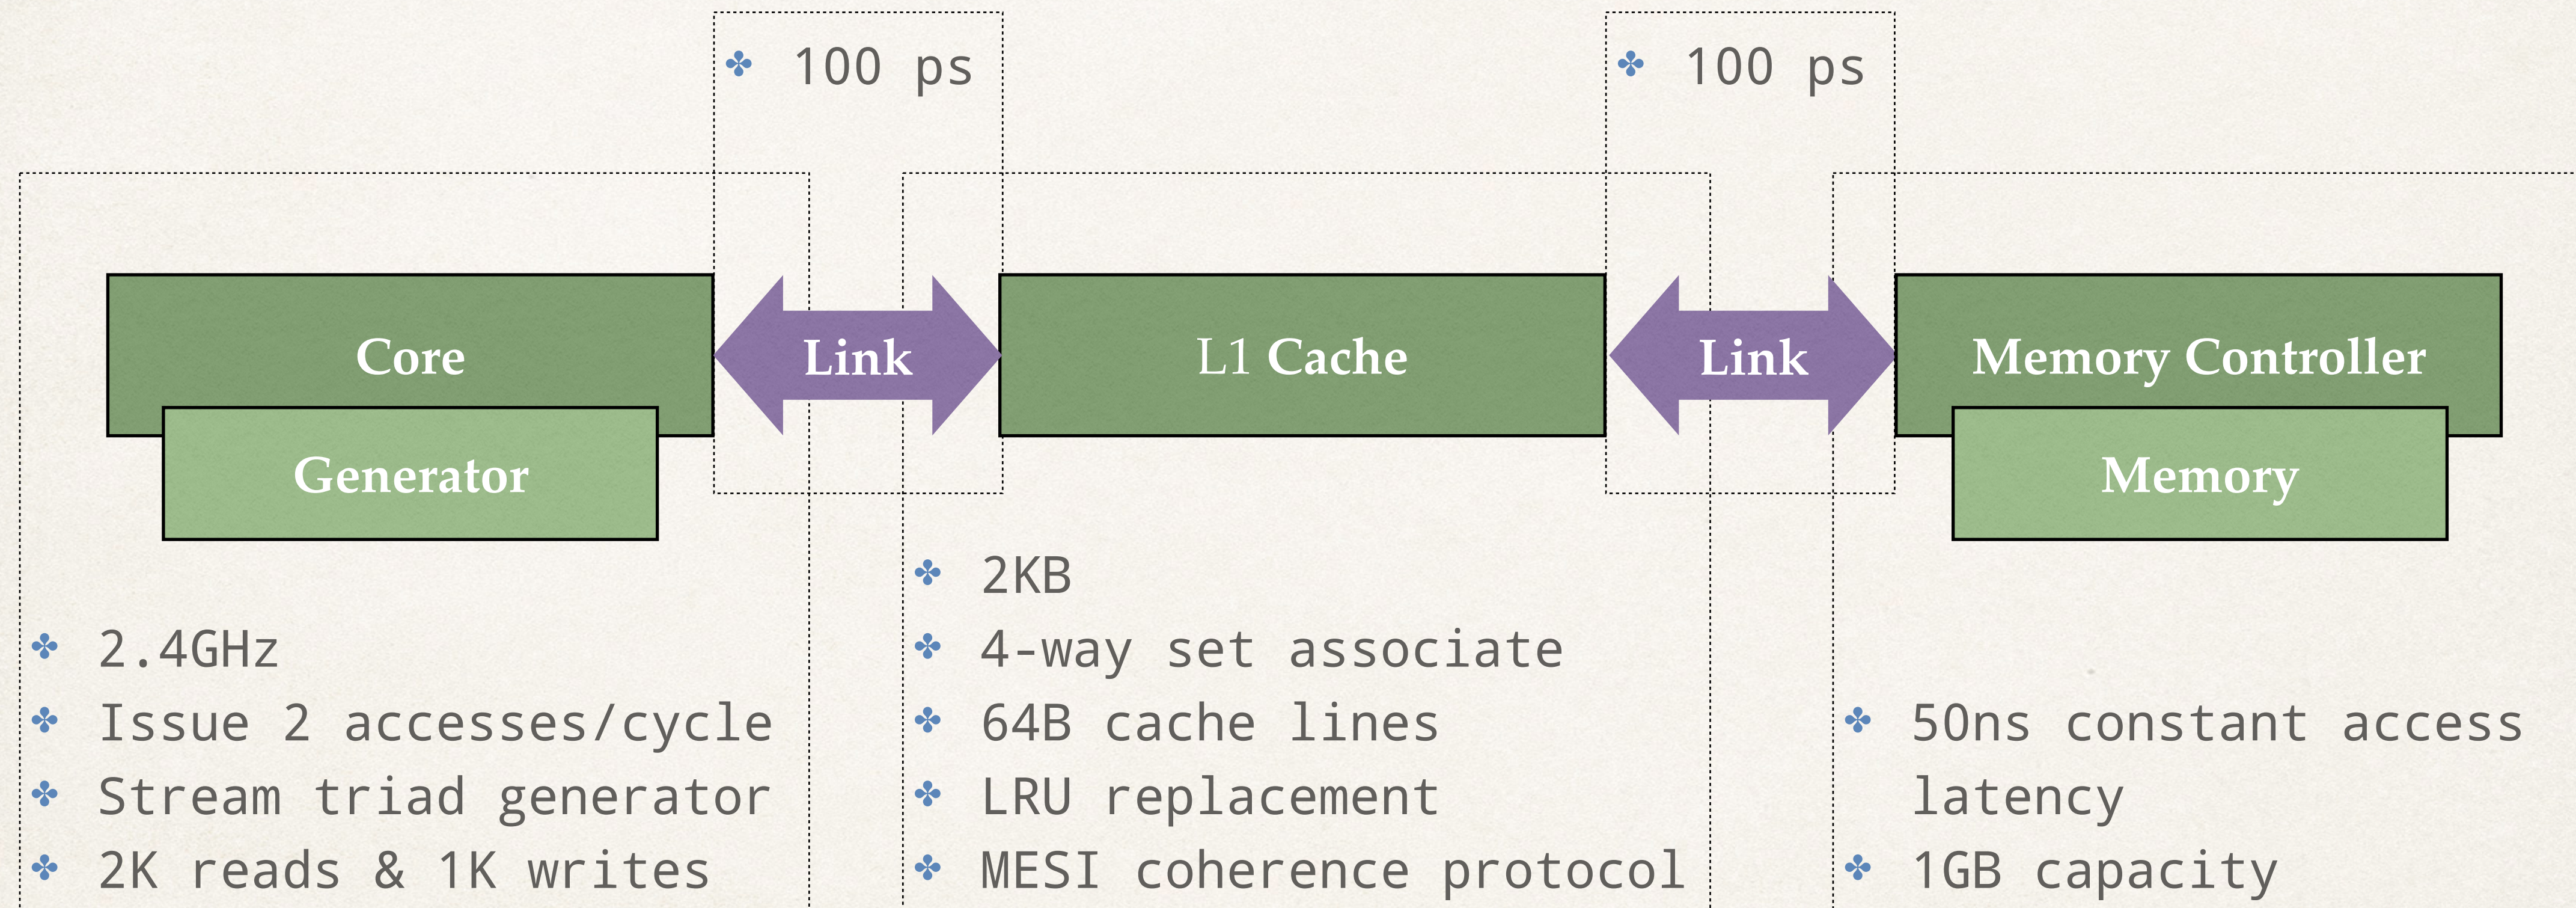
Using ELI version 0.9.0

Compiled on: Apr 13 2022 12:48:52, using file: generators/streambench.h

“sst-info element_name” to extract element and component information.

Stream Triad Kernel:

$a[i] = b[i] + c[i] * \text{SCALAR};$



Stream Triad Kernel:

```
a[i] = b[i] + c[i] * SCALAR;
```


CPU

```
# Create core using Miranda model
core = sst.Component("core", "miranda.BaseCPU")
core.addParams({
    'clock':          "2.4GHz",
    'max_reqs_cycle': 2,
})

# Set and configure a memory access pattern generator on the core
gen = core.setSubComponent("generator", "miranda.STREAMBenchGenerator")
gen.addParams({
    'n':          1000,
    'start_a':    0,          # Start address for array A
    'start_b':    1000 * 32,   # Start address for array B
    'start_c':    2 * 1000 * 32, # Start address for array C
    'operandwidth': 32,       # Size of array elements
})
```

L1 Cache

```
# Ceate Cache using memHierarchy
l1 = sst.Component("L1", "memHierarchy.Cache")
l1.addParams({
    'cache_frequency': clock,
    'cache_size':      "2KB",
    'associativity':    4,
    'access_latency_cycles': 2,
    'L1':              1,
    'cache_line_size': 64,
    'coherence_protocol': "MESI",
    'replacement_policy': "lru",
})
```

Memory

```
# Ceate Memory Controller using memHierarchy
mctrl = sst.Component("memctrl", "memHierarchy.MemController")
mctrl.addParams({
    "clock" : "1GHz",
    "backing" : "none", # Don't allocate space for data values;
                        # Miranda doesn't use them
})

# Use the 'SimpleMem' constant latency memory timing model in our memory
memory = mctrl.setSubComponent("backend", "memHierarchy.simpleMem")
memory.addParams({
    "access_time" : "50ns",
    "mem_size" : "1GiB",
})
```

Link

```
# link0: core <=> cache
link0 = sst.Link("core_to_cache")
link0.connect((core, "cache_link", "300ps"), (l1, "high_network_0", "300ps"))

# link1: cache <=> memory
link1 = sst.Link("cache_to_memory")
link1.connect((l1, "low_network_0", "300ps"), (mctrl, "direct_link", "300ps"))
```


ELEMENT 0 = miranda ()

Num Components = 1

Component 0: BaseCPU

CATEGORY: PROCESSOR COMPONENT

NUM STATISTICS = 17

STATISTIC 0 = split_read_reqs [Number of read requests split over a cache line boundary] (requests) Enable level = 2
STATISTIC 1 = split_write_reqs [Number of write requests split over a cache line boundary] (requests) Enable level = 2
STATISTIC 2 = split_custom_reqs [Number of custom requests split over a cache line boundary] (requests) Enable level = 2
STATISTIC 3 = read_reqs [Number of read requests issued] (requests) Enable level = 1
STATISTIC 4 = write_reqs [Number of write requests issued] (requests) Enable level = 1
STATISTIC 5 = custom_reqs [Number of custom requests issued] (requests) Enable level = 1
STATISTIC 6 = total_bytes_read [Count the total bytes requested by read operations] (bytes) Enable level = 1
STATISTIC 7 = total_bytes_write [Count the total bytes requested by write operations] (bytes) Enable level = 1
STATISTIC 8 = total_bytes_custom [Count the total bytes requested by custom operations] (bytes) Enable level = 1
STATISTIC 9 = req_latency [Running total of all latency for all requests] (ns) Enable level = 2
STATISTIC 10 = cycles_with_issue [Number of cycles which CPU was able to issue requests] (cycles) Enable level = 1
STATISTIC 11 = cycles_no_issue [Number of cycles which CPU was not able to issue requests] (cycles) Enable level = 1
STATISTIC 12 = time [Nanoseconds spent issuing requests] (ns) Enable level = 1
STATISTIC 13 = cycles_hit_fence [Number of issue cycles which stop issue at a fence] (cycles) Enable level = 2
STATISTIC 14 = cycles_max_reorder [Number of issue cycles which hit maximum reorder lookup] (cycles) Enable level = 2
STATISTIC 15 = cycles_max_issue [Cycles with maximum operation issue] (cycles) Enable level = 2
STATISTIC 16 = cycles [Cycles executed] (cycles) Enable level = 1

Most Components and SubComponents define statistics.
\$ sst-info miranda

Set Statistics Level

```
# Enable SST Statistics Outputs for this simulation

# Print statistics of level 5 and below (0-5)
sst.setStatisticLoadLevel(5)
```

Select Statistics of Components

```
# Print statistics of level 5 and below (0-5)
sst.setStatisticLoadLevel(5)

# Enable all statistics for all components
sst.enableAllStatisticsForAllComponents()

# Enable all statistics according component name
sst.enableAllStatisticsForComponentName("L1")

# Enable all statistics according component name and statistic name
sst.enableStatisticsForComponentName("L1", "CacheMisses")
sst.enableStatisticsForComponentName("L1", "CacheHits")
```

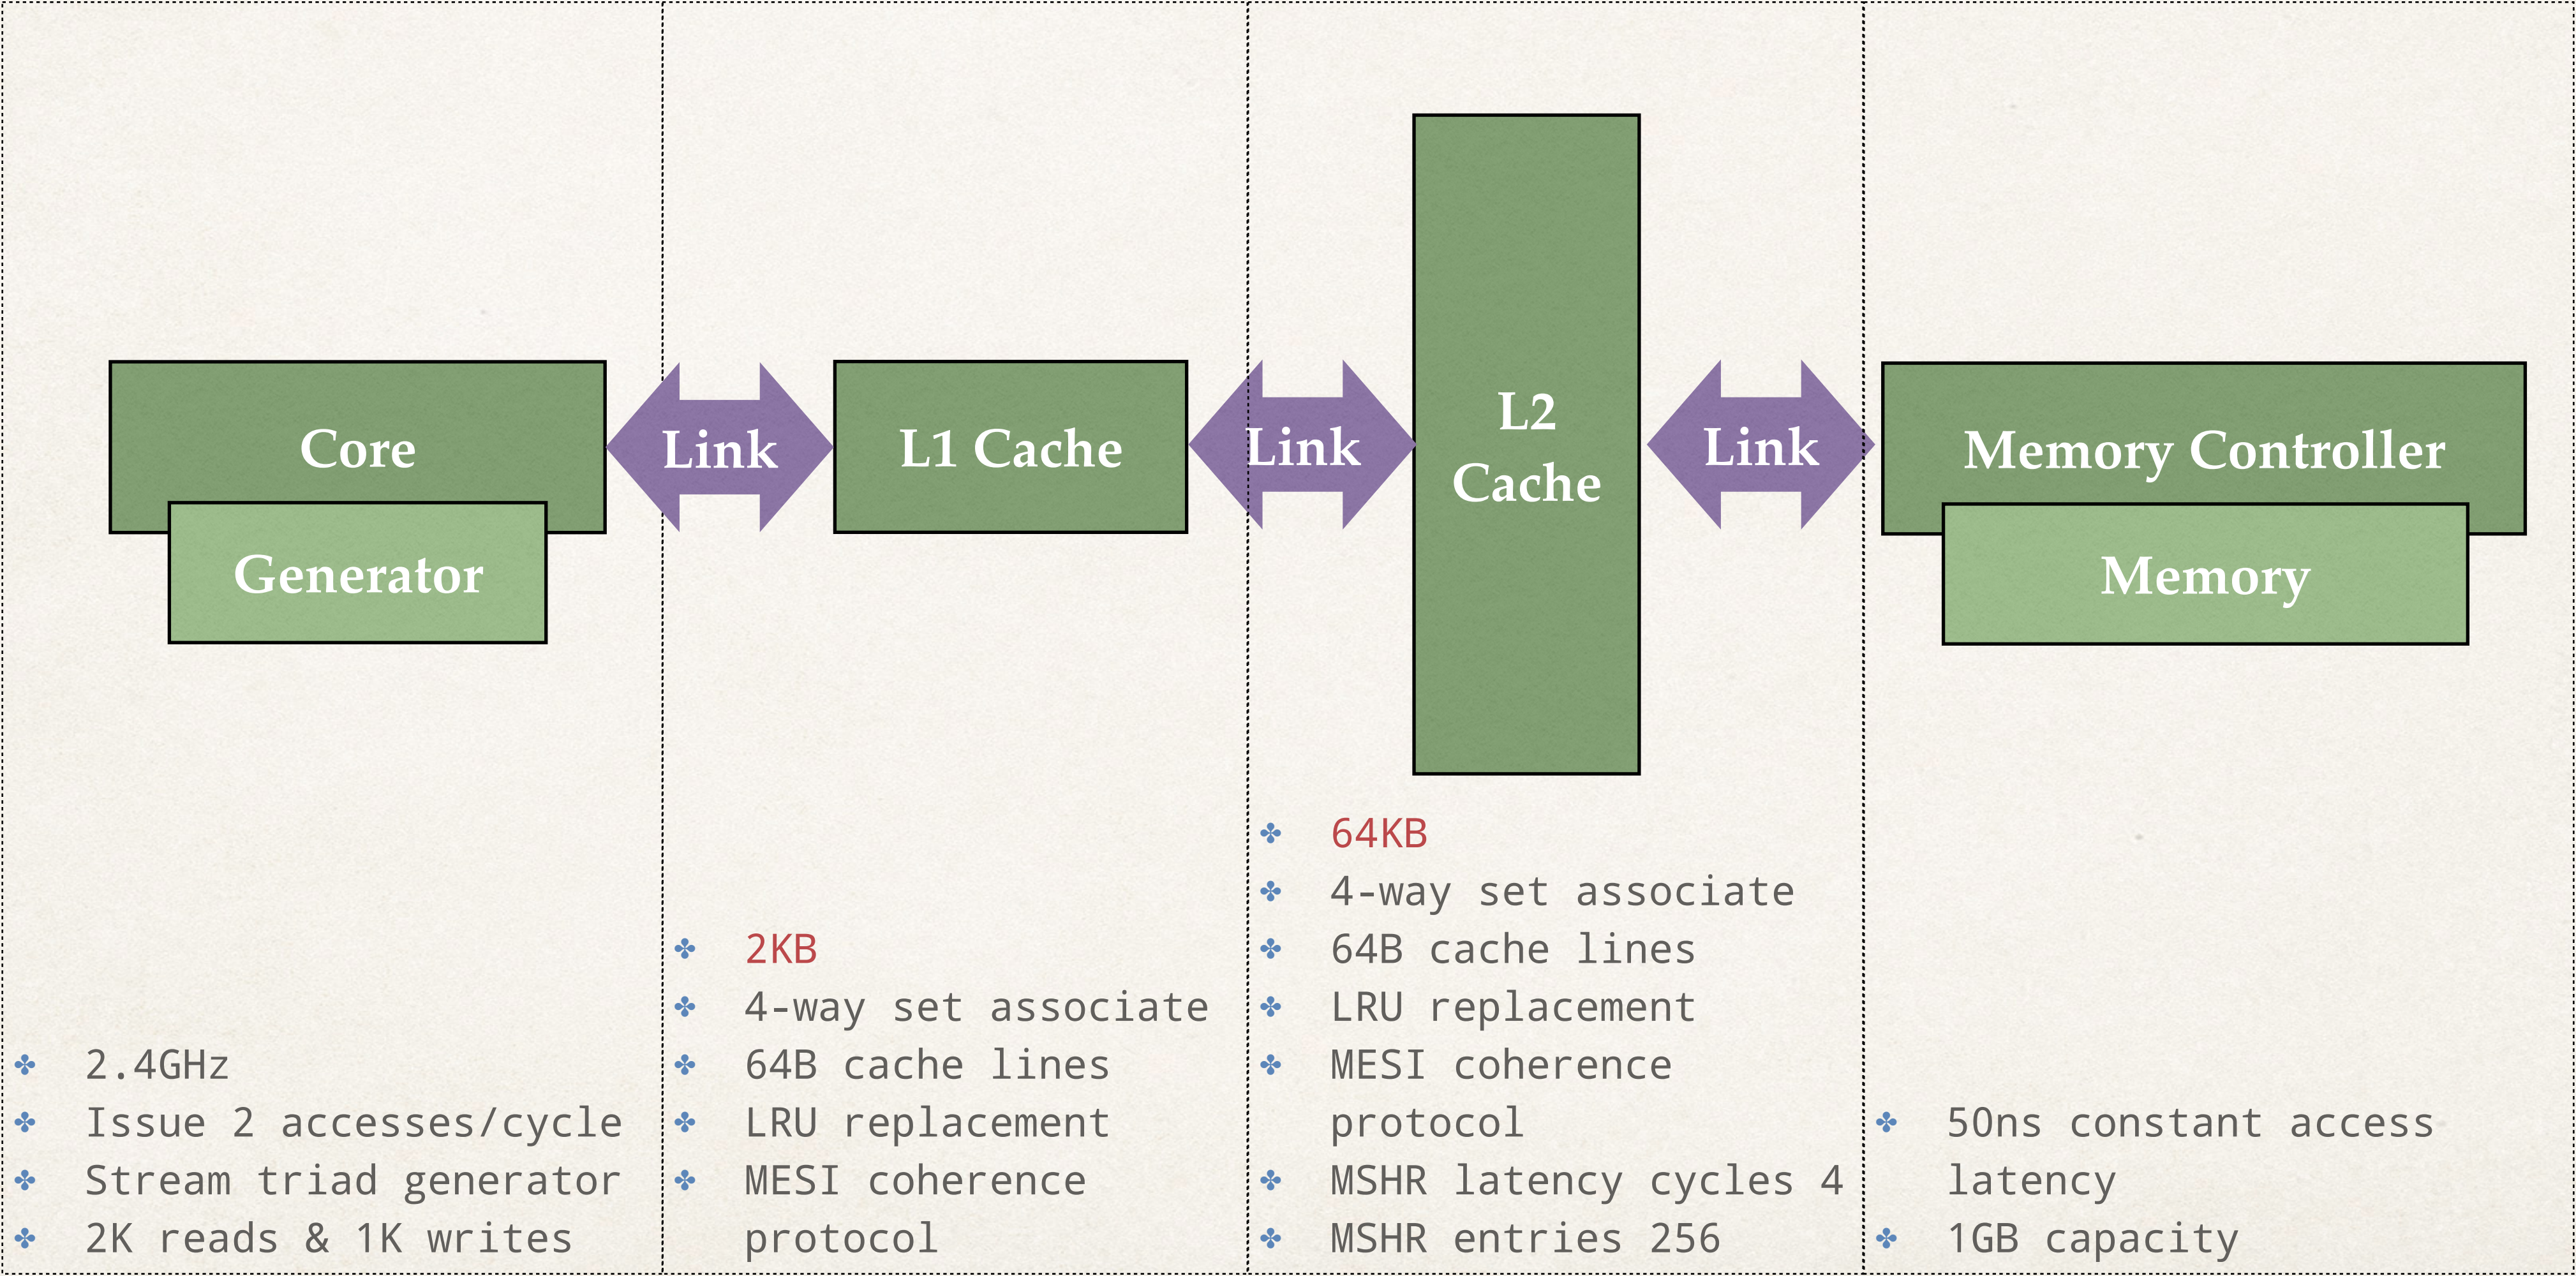
Configure Statistic output

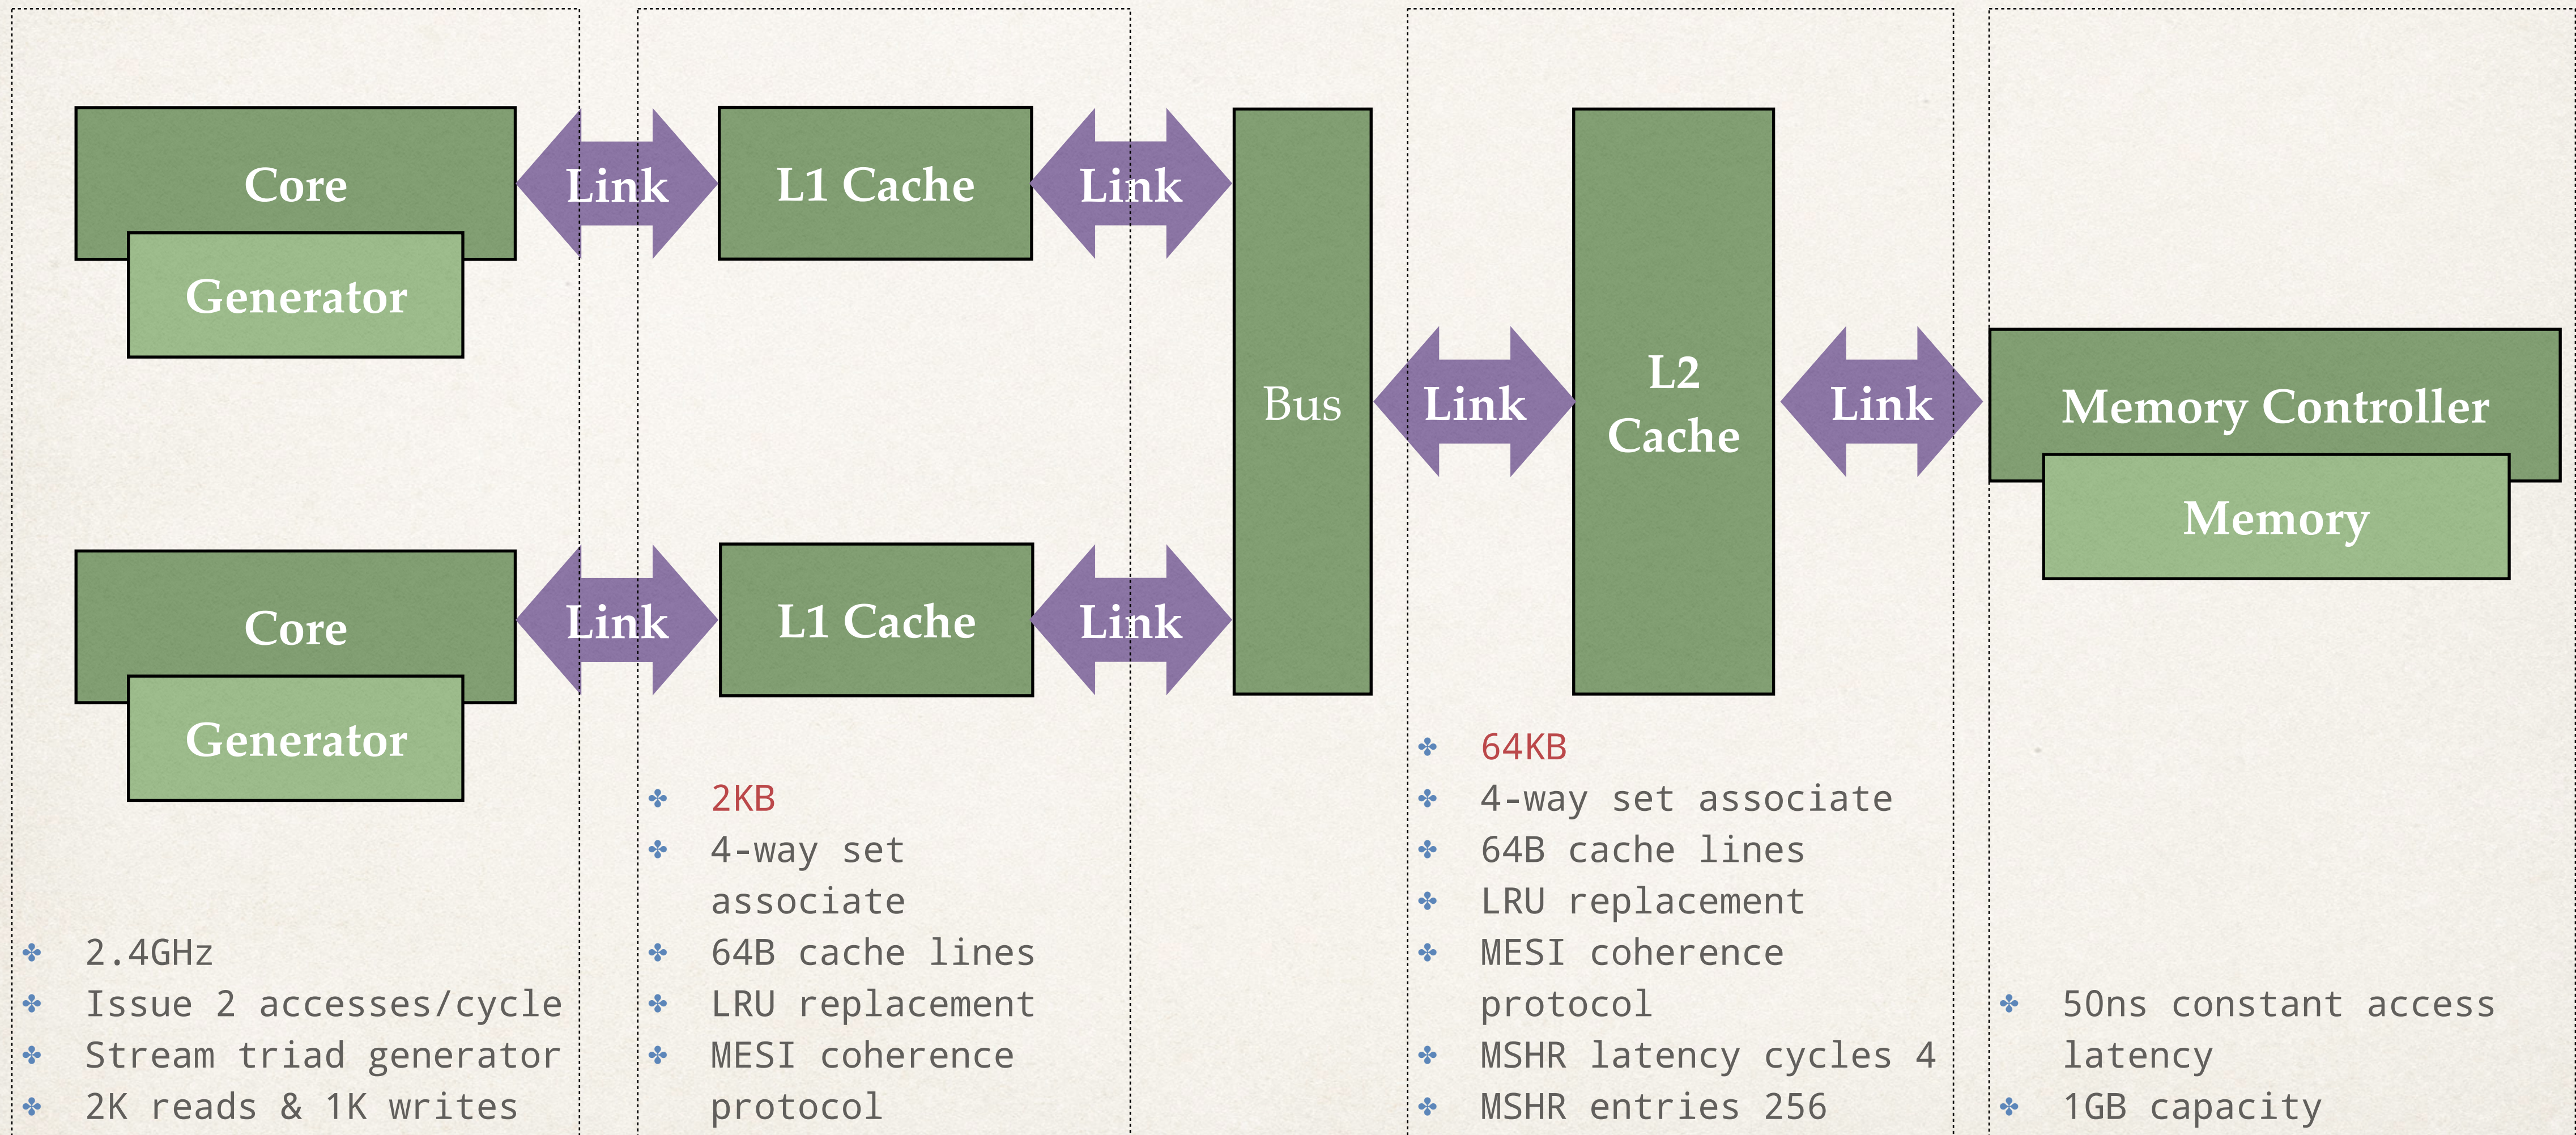
Supported outputs: console, txt, json, hdf5

```
# Generate statistics in CSV format
sst.setStatisticOutput("sst.statoutputcsv")

# Send the statistics to a file called 'stats.csv'
sst.setStatisticOutputOptions( { "filepath" : "stats.csv" })
```


Demo



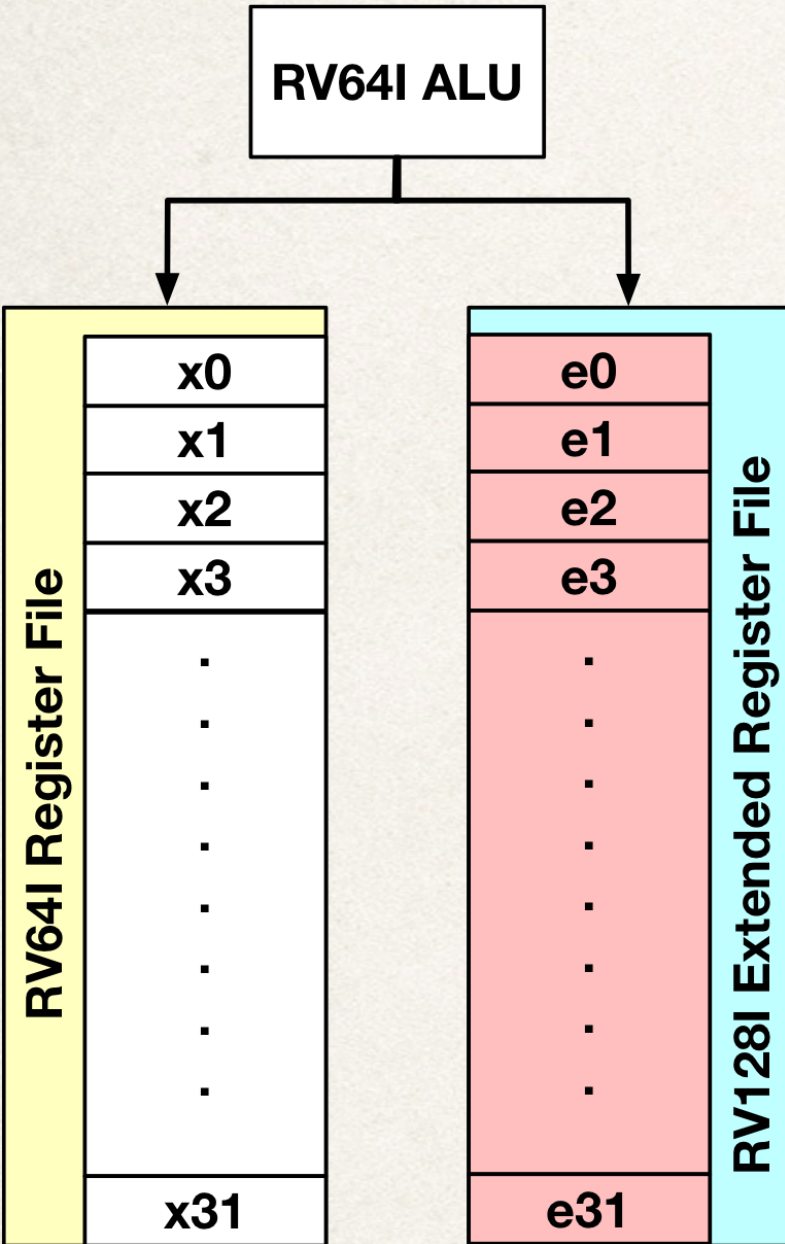


SST for xBGAS

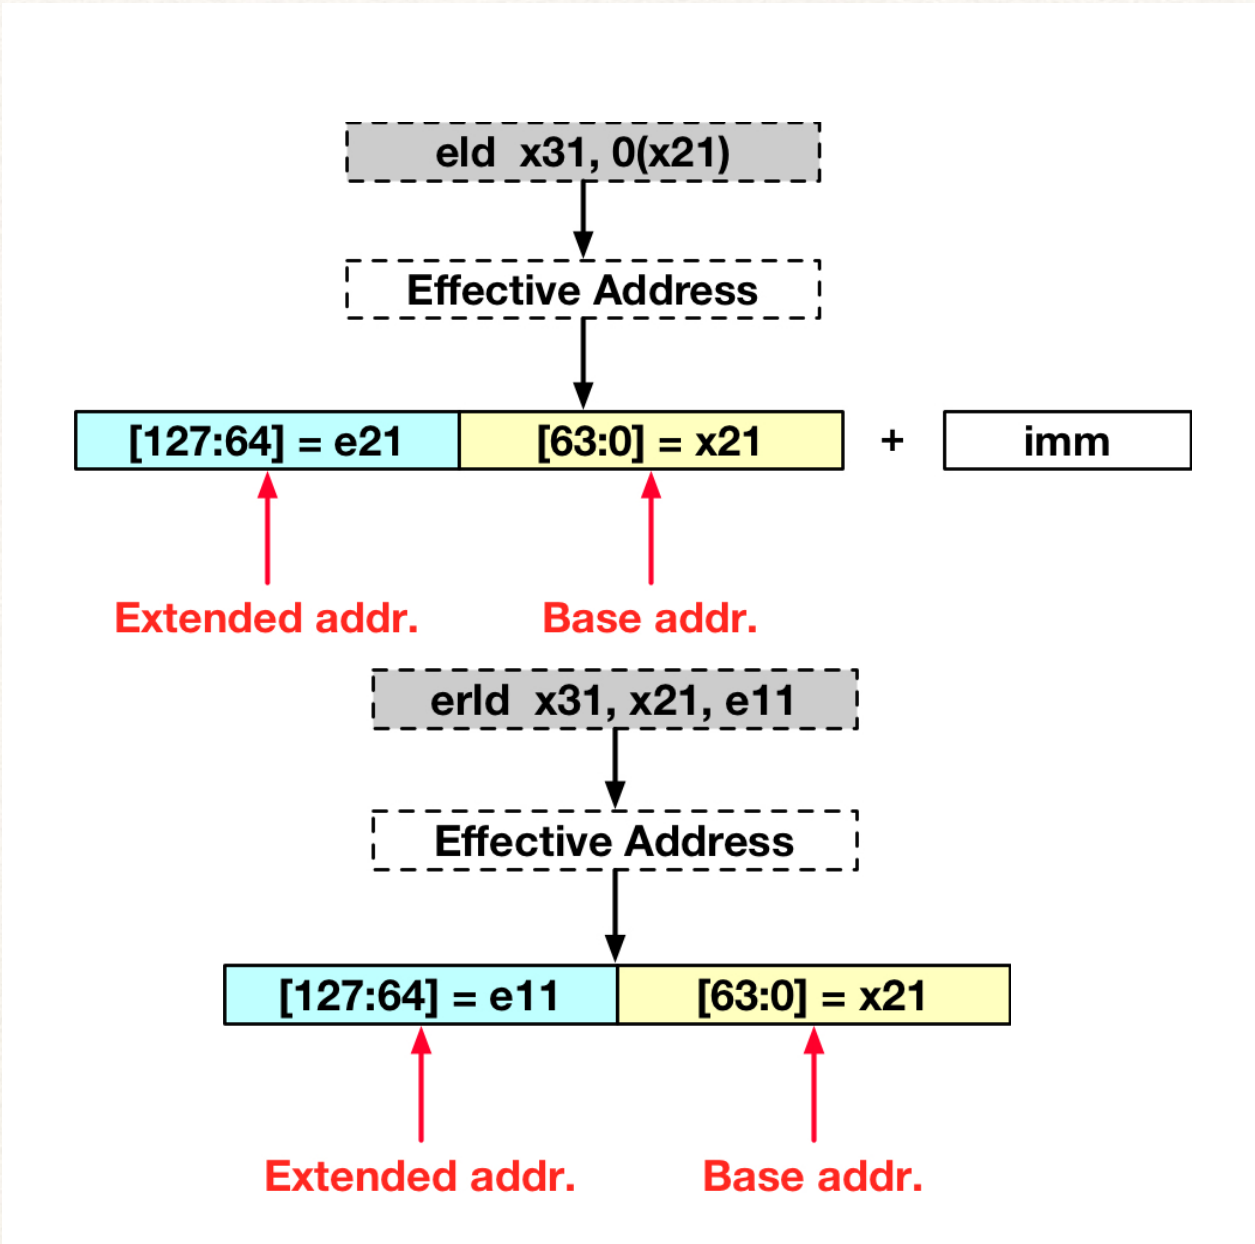
- ❖ Rev SST component:
 - ❖ Designed to provide cycle-based simulation capabilities of **RISC-V** core / cores.
 - ❖ Built upon the standard SST “core” component libraries.
 - ❖ Can be attached to any other existing SST component for full system simulations.
 - ❖ Provides the ability to **load compiled binaries** (ELF binaries).
 - ❖ Supports for **custom instruction extension**.



❖ Register files



Extended Registers



xBGAS Addressing Model

Rev / src / RevInstTable.h

```
namespace SST{
  namespace RevCPU {

    typedef struct{
      uint32_t RV32[_REV_NUM_REGS_];    ///< RevRegFile: RV32I register file
      uint64_t RV64[_REV_NUM_REGS_];    ///< RevRegFile: RV64I register file
      float SPF[_REV_NUM_REGS_];        ///< RevRegFile: RVxxF register file
      double DPF[_REV_NUM_REGS_];       ///< RevRegFile: RVxxD register file

      // Extended register files
      uint32_t ERV32[_REV_NUM_REGS_];   ///< RevRegFile: Extended RV32I register file
      uint64_t ERV64[_REV_NUM_REGS_];   ///< RevRegFile: Extended RV64I register file

      uint32_t RV32_PC;                 ///< RevRegFile: RV32 PC
      uint64_t RV64_PC;                 ///< RevRegFile: RV64 PC
      uint64_t FCSR;                   ///< RevRegFile: FCSR

      uint32_t cost;                    ///< RevRegFile: Cost of the instruction
      bool trigger;                     ///< RevRegFile: Has the instruction been triggered?
      unsigned Entry;                   ///< RevRegFile: Instruction entry
    }RevRegFile;

  }RevProc: register file construct
}
```

Register file struct in RevCPU

❖ Instruction encoding tables

std::vector<RevInstEntry> RV32ETable = {

Rev / src / insns / RV32E.h

// Load operations

{"elw %rd, \$imm(%rs1)", 1, 0b1110111, 0b010, 0b0000000, RegGPR, RegGPR, RegUNKNOWN, RegUNKNOWN, 0b0, FImm, RVTypeI, &elw },

{"elh %rd, \$imm(%rs1)", 1, 0b1110111, 0b001, 0b0000000, RegGPR, RegGPR, RegUNKNOWN, RegUNKNOWN, 0b0, FImm, RVTypeI, &elh },

{"elhu %rd, \$imm(%rs1)", 1, 0b1110111, 0b101, 0b0000000, RegGPR, RegGPR, RegUNKNOWN, RegUNKNOWN, 0b0, FImm, RVTypeI, &elhu },

{"elb %rd, \$imm(%rs1)", 1, 0b1110111, 0b000, 0b0000000, RegGPR, RegGPR, RegUNKNOWN, RegUNKNOWN, 0b0, FImm, RVTypeI, &elb },

{"elbu %rd, \$imm(%rs1)", 1, 0b1110111, 0b100, 0b0000000, RegGPR, RegGPR, RegUNKNOWN, RegUNKNOWN, 0b0, FImm, RVTypeI, &elbu },

// Store Operations

{"esw %rs2, \$imm(%rs1)", 1, 0b1111011, 0b010, 0b0000000, RegIMM, RegGPR, RegGPR, RegUNKNOWN, 0b0, FImm, RVTypeS, &esw },

{"esh %rs2, \$imm(%rs1)", 1, 0b1111011, 0b001, 0b0000000, RegIMM, RegGPR, RegGPR, RegUNKNOWN, 0b0, FImm, RVTypeS, &esh },

{"esb %rs2, \$imm(%rs1)", 1, 0b1111011, 0b000, 0b0000000, RegIMM, RegGPR, RegGPR, RegUNKNOWN, 0b0, FImm, RVTypeS, &esb },

Mnemonic, cost, opcode, funct3, funct7, rdClass, rs1Class, rs2Class, rs3Class, imm12, imm, format, func

Define the register class for the corresponding slot.
Possible values: RegGPR, RegCSR, RegIMM, RegUNKNOWN
RegUNKNOWN if the field is not utilized.

❖ Instruction function implementation

```
static bool lw(RevFeature *F, RevRegFile *R, RevMem *M, RevInst Inst) {
    if( F->IsRV32() ){
        SEXT(R->RV32[Inst.rd], M->ReadU32( (uint64_t)(R->RV32[Inst.rs1] + (int32_t)(td_u32(Inst.imm, 12))), 32);
        R->RV32_PC += Inst.instSize;
    }else{
        SEXT(R->RV64[Inst.rd], M->ReadU32( (uint64_t)(R->RV64[Inst.rs1] + (int32_t)(td_u32(Inst.imm, 12))), 64);
        R->RV64_PC += Inst.instSize;
    }
    // update the cost
    R->cost += M->RandCost(F->GetMinCost(), F->GetMaxCost());
    return true;
}
```

Standard RV32I

```
static bool elw(RevFeature *F, RevRegFile *R, RevMem *M, RevNIC *N, RevInst Inst) {
    if( F->IsRV32() ){
        // Add NIC object that reads the remote memory
        SEXT( R->RV32[Inst.rd], N->ReadU32( R->ERV32, R->RV32[Inst.rs1] ) + (int32_t)(td_u32(Inst.imm, 12)), 32);
        R->RV32_PC += Inst.instSize;
    }else{
        SEXT( R->RV64[Inst.rd], N->ReadU32( R->ERV64, R->RV64[Inst.rs1] ) + (int32_t)(td_u32(Inst.imm, 12)), 64);
        R->RV64_PC += Inst.instSize;
    }
    // update the cost
    R->cost += N->RandCost(F->GetMinCost(), F->GetMaxCost());
    return true;
}
```

xBGAS extension

- ❖ Utilize the **RevMem** object to read the desired number of bytes via the ReadU32 routine.
- ❖ Utilize the **RevNIC** object to read the desired number of bytes via the ReadU32 routine.
- ❖ Read routines are not exist in RevNIC.
- ❖ NLB (Namespace Lookaside Buffer)
- ❖ Initialization of PE id and its corresponding address

-
- ❖ Presented the design of SST and its key objects.
 - ❖ Simulated a multi-level memory system with three different configurations.
 - ❖ Discussed the approaches to modify the RevCPU model for supporting xBGAS simulation.
 - ❖ Will dive deep into the RevNIC component to understand how it works.
 - ❖ Will implement the remote memory operation routines to support the instruction functions.
 - ❖ SST is not limited to simulating HPC Systems; SST GitHub repo provides an example of carwash simulation.
 - ❖ <https://github.com/sstsimulator/sst-elements/tree/master/src/sst/elements/simpleSimulation>
 - ❖ <https://www.youtube.com/watch?v=ipxoVb2dEcw&list=PLgehegDe4T2y1badxrx cuvIsX42V64t2x>

Thank you!

Q&A